CARNEGIE MELLON UNIVERSITY
COMPUTER SCIENCE DEPARTMENT
15-445/645 – DATABASE SYSTEMS (FALL 2020)
PROF. ANDY PAVLO

Homework #5 (by Preetansh Goyal)
Due: **Tuesday Dec 6, 2020 @ 11:59pm**

IMPORTANT:
- **Upload this PDF** with your answers to **Gradescope by 11:59pm on Tuesday Dec 6, 2020**.
- **Plagiarism**: Homework may be discussed with other students, but all homework is to be completed **individually**.
- **You have to use this PDF for all of your answers.**

For your information:
- Graded out of **86** points; **3** questions total

Revision : 2020/12/07 22:14

Question	Points	Score
Write-Ahead Logging	35	
Replication	33	
Miscellaneous	18	
Total:	86	

Question 1: Write-Ahead Logging . . . . [35 points]

Consider a DBMS using write-ahead logging with physical log records with the STEAL and NO-FORCE buffer pool management policy. Assume the DBMS executes a non-fuzzy checkpoint where all dirty pages are written to disk.

Its transaction recovery log contains log records of the following form:

    <txnId, objectId, beforeValue, afterValue>

The log also contains checkpoint, transaction begin, and transaction commit records.

The database contains three objects (i.e., X, Y, and Z).

The DBMS sees records as in Figure 1 in the WAL on disk after a crash.

LSN	WAL Record
1	<T1 BEGIN>
2	<T1, X, 1, 2>
3	<T2 BEGIN>
4	<T2, Y, 1, 2>
5	<T1 COMMIT>
6	<T2, Y, 2, 3>
7	<T3 BEGIN>
8	<T3, Z, 1, 2>
9	<T2, X, 2, 3>
10	<**CHECKPOINT**>
11	<T2, Y, 3, 4>
12	<T3, Z, 2, 3>
13	<T3 COMMIT>
14	<T2, Z, 3, 4>

Figure 1: WAL

(a) **[10 points]** What are the values of X, Y, and Z in the database stored on disk before the DBMS recovers the state of the database?

    □ X=1, Y=1, Z=1

    □ X=2, Y=1, Z=1

    □ X=3, Y=4, Z=4

    □ X=3, Y=4, Z=3

    □ X=2, Y:4, Z=Not possible to determine

    □ X=3, Y:Not possible to determine, Z=4

    □ X=3, Y,Z:Not possible to determine

    □ X=2, Y:Not possible to determine, Z=4

    □ X=2, Y,Z:Not possible to determine

    □ X,Y,Z:Not possible to determine

(b) **[5 points]** What should be the correct action on T1 when recovering the database from WAL?

    ☐ undo all of T1's changes

    ☐ redo all of T1's changes

    ☐ do nothing to T1

(c) **[5 points]** What should be the correct action on T2 when recovering the database from WAL?

    ☐ undo all of T2's changes

    ☐ redo all of T2's changes

    ☐ do nothing to T2

(d) **[5 points]** What should be the correct action on T3 when recovering the database from WAL?

    ☐ undo all of T3's changes

    ☐ redo all of T3's changes

    ☐ do nothing to T3

(e) **[10 points]** Assume that the DBMS flushes all dirty pages when the recovery process finishes. What are the values of X, Y, and Z after the DBMS recovers the state of the database from the WAL in Figure 1?

    ☐ X=1, Y=1, Z=1

    ☐ X=1, Y=2, Z=3

    ☐ X=2, Y=1, Z=3

    ☐ X=2, Y=1, Z=4

    ☐ X=2, Y=3, Z=2

    ☐ X=2, Y=4, Z=4

    ☐ X=3, Y=4, Z=3

    ☐ Not possible to determine

Question 2: Replication . . . . [33 points]

Consider a DBMS using active-passive, master-replica replication with multi-versioned concurrency control. All read-write transactions go to the master node (NODE A), while read-only transactions are routed to the replica (NODE B). You can assume that the DBMS has "instant" fail-over and master elections. That is, there is no time gap between when the master goes down and when the replica gets promoted as the new master. For example, if NODE A goes down at timestamp ① then NODE B will be elected the new master at ②. Note that this is not a realistic assumption but we're using it to simplify the problem setup.

The database has a single table foo(id,val) with the following tuples:

id	val
1	a
2	b
3	c

Table 1: **foo(id,val)**

For each questions listed below, assume that the following transactions shown in Figure 2 are executing in the DBMS: (1) Transaction #1 on NODE A and (2) Transaction #2 on NODE B. You can assume that the timestamps for each operation is the real physical time of when it was invoked at the DBMS and that the clocks on both nodes are perfectly synchronized (again, this is not a realistic assumption).

time	operation
①	BEGIN;
②	UPDATE foo SET val = 'aa';
③	UPDATE foo SET val = 'aaa' WHERE id = 1;
④	UPDATE foo SET val = 'bb' WHERE id = 2;
⑤	COMMIT;

(a) Transaction #1 – NODE A

time	operation
②	BEGIN READ ONLY;
③	SELECT val FROM foo WHERE id = 1;
④	SELECT val FROM foo WHERE id = 2;
⑤	SELECT val FROM foo WHERE id = 2;
⑥	COMMIT;

(b) Transaction #2 – NODE B

Figure 2: Transactions executing in the DBMS.

(a) Assume that the DBMS is using *asynchronous* replication with *continuous* log streaming (i.e., the master node sends log records to the replica in the background after the transaction executes them). Suppose that NODE A crashes at timestamp ⑤ <u>before</u> it executes the COMMIT operation.

    i. **[10 points]** If Transaction #2 is running under READ COMMITTED, what is the return result of the val attribute for its SELECT query at timestamp ④? Select all that are possible.

       ☐ a

       ☐ aa

       ☐ aaa

       ☐ b

☐ bb
☐ c
☐ None of the above

ii. **[10 points]** If Transaction #2 is running under the READ UNCOMMITTED isolation level, what is the return result of the val attribute for its SELECT query at timestamp ⑤? Select all that are possible.

☐ a
☐ aa
☐ aaa
☐ b
☐ bb
☐ c
☐ None of the above

(b) **[13 points]** Assume that the DBMS is using *asynchronous* replication with *on commit* propagation. Suppose that both NODE A and NODE B crash at exactly the same time at timestamp ⑥ <u>after</u> executing Transaction #1's COMMIT operation. You can assume that the application was notified that the Transaction #1 was committed successfully.

After the crash, you find that NODE A had a major hardware failure and cannot boot. NODE B is able to recover and is elected the new master.

What are the values of the tuples in the database when the system comes back online? Select all that are possible.

☐ { (1,a), (2,b), (3,c) }
☐ { (1,aa), (2,aa), (3,aa) }
☐ { (1,aaa), (2,bb), (3,c) }
☐ { (1,aaa), (2,bb), (3,aa) }
☐ { (1,a), (2,bb), (3,c) }
☐ { (1,a), (2,bb), (3,aa) }
☐ None of the above

Question 3: Miscellaneous . . . . [18 points]

(a) **[6 points]** Under the NO-STEAL policy, a DBMS will need to store the whole table in RAM if a transaction updates all the records of that table.

- ☐ True
- ☐ False

(b) **[6 points]** In ARIES, log records are immediately flushed on the log, as soon as they are produced.

- ☐ True
- ☐ False

(c) **[6 points]** Without checkpoints, the redo phase of ARIES recovery should process the whole log.

- ☐ True
- ☐ False